

API Health Monitoring System — DevOps Internship Assignment 2026

1) Problem Statement

Teams often rely on a few critical APIs and need to know quickly when one goes down. This project builds a self-hosted monitoring system that checks API endpoints at a fixed interval, detects meaningful state changes, and notifies users—without using managed monitoring services. The goal is not to feature completeness but clear system design, reliability thinking, and solid infrastructure choices.

2) High-Level Architecture

The system is intentionally simple and AWS-native:

EC2 instance runs the Python monitoring app on a cron schedule.

DynamoDB stores:

- API configuration (api_health_configs)
- Last known health state (api_health_states)
- SNS delivers email alerts on state changes.
- IAM role on EC2 provides least-privilege access to DynamoDB + SNS.

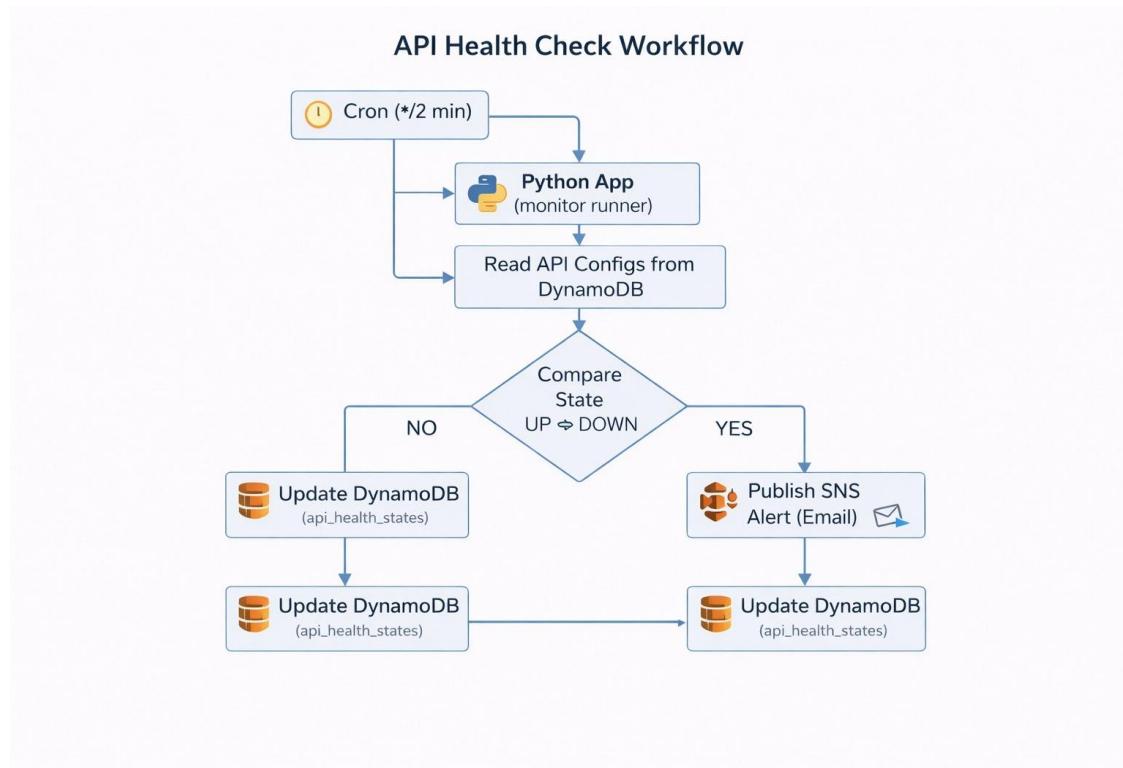


Fig1.API Health Check Work-flow

3) Component Interactions & Data Flow

Health Check Flow (per run)

- Cron triggers the Python app every 2 minutes.
- App reads all API configs from DynamoDB.
- Each API is checked using requests with timeout + expected status rules.
- The result is compared with the last stored state.
- If the state changed (UP ↔ DOWN), an SNS alert is published.
- The new state is saved back to DynamoDB.

4) Infrastructure & Deployment (Terraform)

Infrastructure is defined in infra/terraform:

- EC2: runs the cron + Python app
- DynamoDB: api_health_configs, api_health_states
- SNS: alert topic
- IAM: least-privilege policy attached to EC2 role
- Outputs: EC2 public IP, table names, SNS topic ARN

Deployment Steps:

1. `terraform init / plan / apply`
2. SSH to the EC2 instance
3. `pip install -r app/requirements.txt`

Set env vars:

- `AWS_REGION, APP_ENVIRONMENT`
- `DDB_CONFIG_TABLE, DDB_STATE_TABLE`
- `SNS_TOPIC_ARN`
- Configure cron to run `python3 -m app.src.cron_entrypoint`

5) Code Structure

- `app/src`
- `config.py`: loads env config
- `models.py`: dataclasses for config + state
- `dynamodb_client.py`: scan configs, get/update states
- `sns_client.py`: publish alerts
- `health_checker.py`: HTTP request + status/latency/error logic
- `monitor_runner.py`: orchestrates one run
- `cron_entrypoint.py`: cron entrypoint

- infra/terraform
- dynamodb.tf, sns.tf, iam.tf, ec2.tf, outputs.tf, etc.

6) Design Decisions, Assumptions, and Trade-offs

- Scheduling: Cron on EC2 is simple and transparent.
- Trade-off: if the instance is down, checks pause.
- Storage: DynamoDB on-demand avoids capacity planning.
- Trade-off: a full table scan doesn't scale forever.
- Alerting: SNS email is straightforward and cheap.
- Trade-off: limited customization.
- Logging: cron redirects to a local log file.
- Trade-off: not centralized.
- Assumptions: moderate number of APIs, per-minute checks, single AWS region.

7) Scalability & Reliability Considerations

- DynamoDB scales automatically for growth.
- Multiple EC2 workers could split work for larger fleets.
- Future scale could use a queue (e.g., SQS) with worker pools.
- State is stored persistently, so failures don't lose history.

8) Validation Steps

- Insert a test API config
- Run `python3 -m app.src.cron_entrypoint` once
- Verify a new record in `api_health_states`
- Flip to a failing URL and re-run to confirm SNS email alert

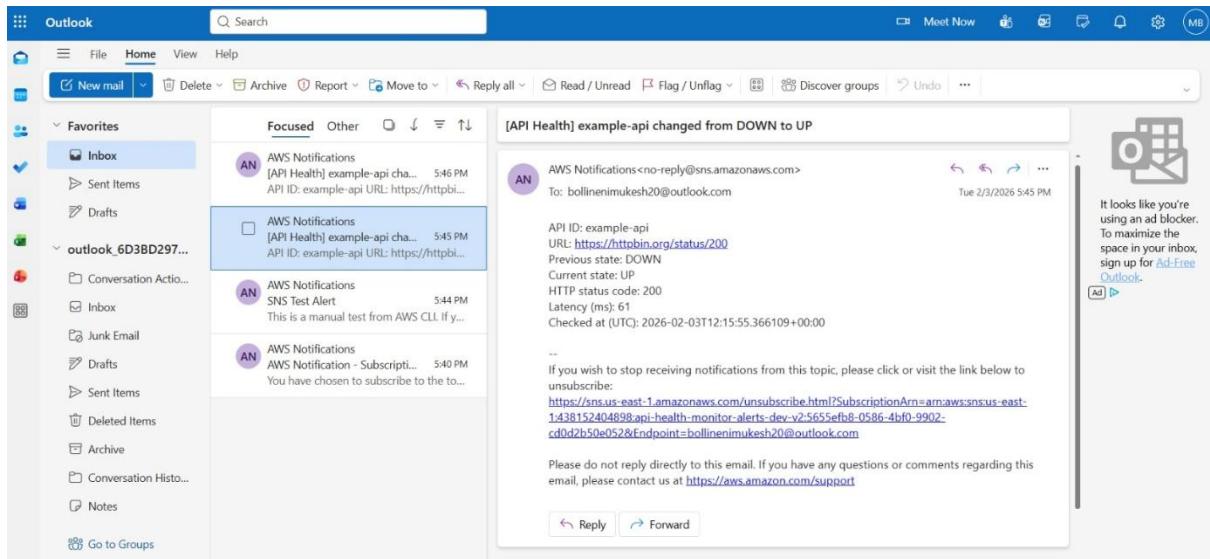


Fig2. API up email

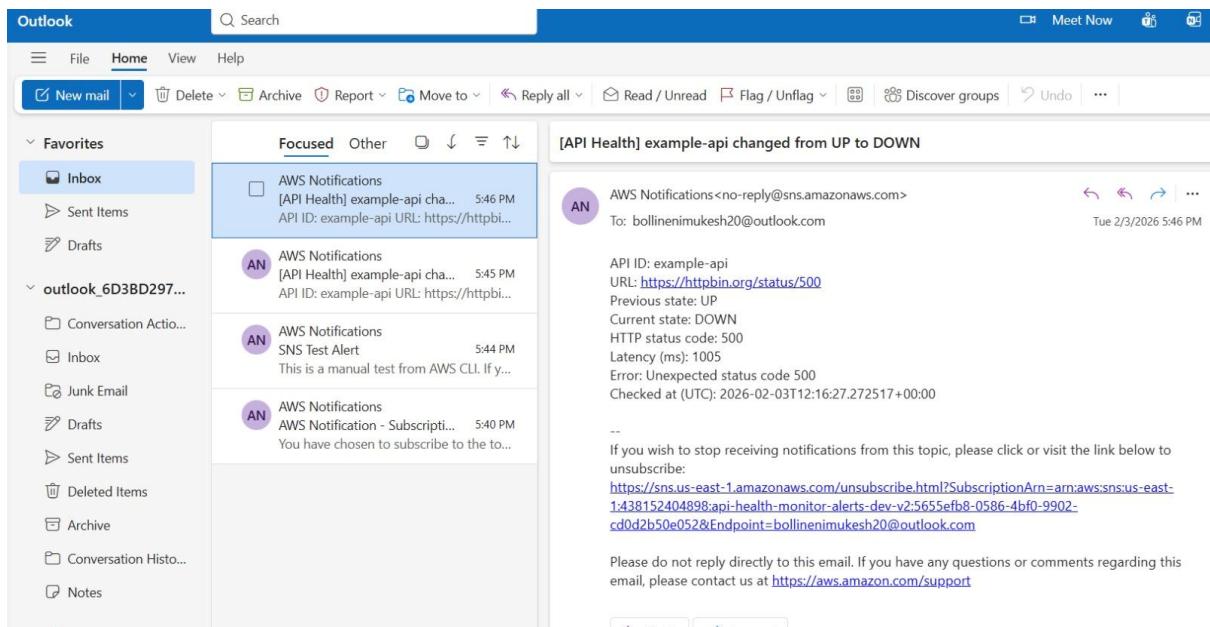


Fig3. API down email

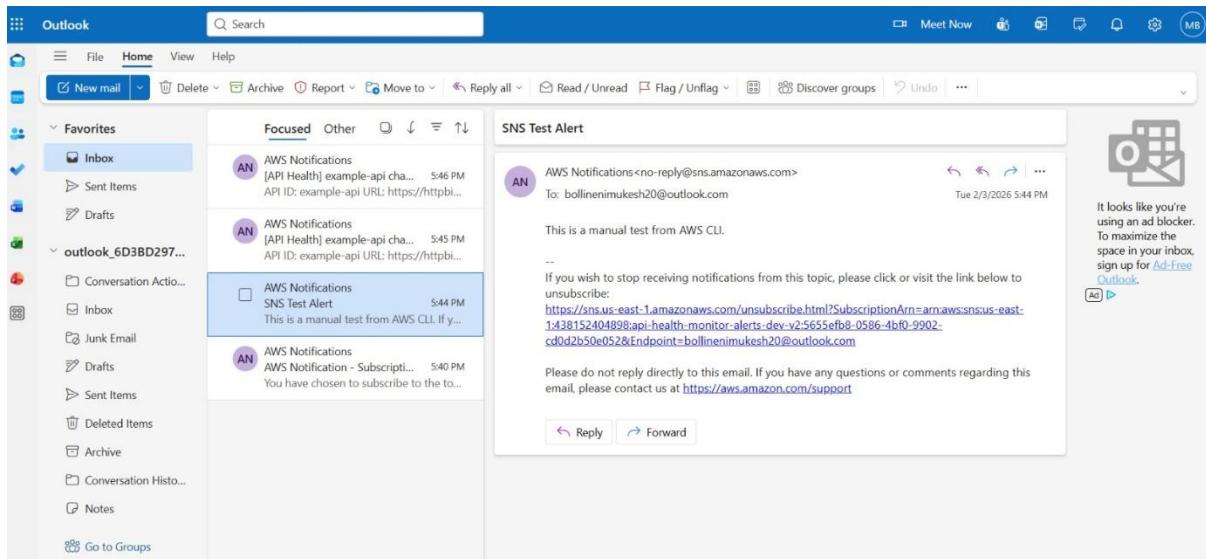


Fig4. API SNS alert mail

9) Diagram

The attached diagram shows:

High-level architecture

Data/alert flow

AWS logical layout




Fig5. Self-Hosted API Health Monitoring System – Architecture and Workflow